

Project Documentation Structure (2025 Best Practices)

Below is a modular documentation outline for an AI-enabled web app, following modern standards. Each section is presented as a separate markdown file with best-practice content, diagrams, and citations to authoritative sources.

README.md

This README introduces the project and provides quick-start information. It should *orient new users* by summarizing goals, stack, and usage ¹.

- **Project Name & Purpose:** A clear title and one-line description of what the app does (e.g. "AI Helper App – a LangChain-based assistant with OpenAI and vector search").
- **Summary:** A concise summary explaining the problem domain, stakeholders, and high-level design. Include goals and target users. Use domain-specific terms for context ².
- **Tech Stack:** List major technologies (e.g. Next.js, Cloudflare Workers, LangChain, OpenAI API, Pinecone, Notion API, Ollama) with links. This helps readers quickly grasp dependencies ³.
- **Quick Start / Setup:** Step-by-step bullets for local setup:
 - Prerequisites (Node.js, Wrangler CLI, API keys).
 - Clone repo, install dependencies (e.g. `npm install`).
 - Environment configuration: reference a `.env` or config file for secrets (do NOT commit secrets ⁴).
 - How to run locally (e.g. `npm run dev` for Next.js, `wrangler dev` for Workers).
- **Usage:** Example commands or scenarios (e.g. hitting `/api/chat` route, using the web UI). Include sample requests/responses or screenshots.
- **Contributing:** Brief note pointing to a CONTRIBUTING.md (if any). Encourage issues/PRs, mention code style checks (e.g. linting rules). List core maintainers or teams. A good README often includes contributor credits or references ⁵.
- **Links:** Pointers to other documentation pages (Architecture, Security, etc.). This creates a navigable docs tree.

Example README sections are structured as above to align with community best practices ¹ ².

ARCHITECTURE.md

This document details system design, components, and deployment. It uses Mermaid diagrams to illustrate the architecture.

Figure: Example full-stack architecture on Cloudflare

The core architecture uses Cloudflare's edge network for deployment. A user's request (Client) hits Cloudflare's global edge (with WAF, CDN) and is routed to the Next.js app running as a Cloudflare Pages/Workers application. Backend logic (AI/chat handlers) runs on Cloudflare Workers or Edge Functions.

Persistent storage (e.g. Pinecone vector DB, Notion API) and external LLMs (OpenAI, local Ollama) are accessed via secure APIs. See the diagram above for an overview of traffic flow and services.

```
graph LR
  subgraph Browser
    U(User)
  end
  subgraph Cloudflare
    CF(Edge (Pages/Workers))
  end
  subgraph Backend
    API[Next.js API Routes<br/>(OpenAI, Pinecone, Notion, Ollama)]
    OpenAI[OpenAI API]
    Pinecone[Pinecone Vector DB]
    NotionAPI[Notion API]
    Ollama(Ollama LLM)
  end
  U --> CF
  CF --> API
  API --> OpenAI
  API --> Pinecone
  API --> NotionAPI
  API --> Ollama
```

- **Component Interaction:** The Next.js route handlers (in `pages/api/`) call external services. For example, the `/chat` handler sends prompts to OpenAI or Ollama, stores/retrieves vectors from Pinecone, and reads/writes data in Notion.
- **Groups and Services:** In Mermaid, we use *groups* and *services* to show architecture hierarchies ⁶. Cloudflare's edge (blue box) contains the Next.js service, which connects to external services.
- **Deployment:** Source code is built in GitHub Actions, producing a static frontend (if any) and deploying Workers via Wrangler. Infrastructure is fully serverless. Secrets (API tokens) are stored in Cloudflare's Secret Store and GitHub Secrets ⁷.

Mermaid and reference docs like Cloudflare's confirm that architecture diagrams can clearly show service relationships ⁶ ⁸.

LANGCHAIN_AGENTS.md

This file covers the LangChain agent setup, tool registry, and error handling patterns.

Figure: LangChain multi-agent architecture (Planner, Executor, Communicator, Evaluator) ⁹ ¹⁰

- **Agent Architecture:** Our system uses a LangChain *Planner/Executor* style agent. A high-level *Planner* agent breaks user tasks into subtasks; multiple *Executor* agents (e.g. for RAG, code execution, knowledge retrieval) act on them. A *Communicator* or *Assistant* model interacts with the user, and an *Evaluator* agent may verify outputs. This multi-agent setup (illustrated above) is a known pattern ⁹ ¹⁰.
- **Tool Registry:** We register tools that agents can use, such as database access, external APIs, or custom functions. Optionally, we integrate a Smithery/MCP tool registry: a searchable service

(MCP servers) where models can discover tools dynamically ¹¹ ¹² . This enables extensibility by querying an external registry.

- **Fallback Strategies:** LangChain supports *fallbacks* when an LLM or chain fails ¹³ ¹⁴ . For instance, if OpenAI's chat API times out or hits a rate limit, we can automatically retry with a different model (like Ollama or an alternate prompt) using `withFallbacks()` . LangChain examples show falling back from GPT-3.5 to a GPT-4 model on error ¹⁴ , or to a local Ollama when cloud APIs fail. This ensures resilience against downtime or rate limits.
- **Error Handling:** We employ try/catch around tool calls and logging callbacks (as LangChain suggests) ¹⁵ . Key patterns include:
 - *Retry Logic:* For transient errors (network issues or rate limits), we retry calls with backoff.
 - *Logging:* All errors and tool usage are logged via LangChain callbacks for observability.
 - *Custom Exceptions:* We throw clear exceptions if tools return invalid data, enabling graceful recovery.
 - *User Feedback:* In the chain, we catch exceptions to return user-friendly messages (e.g. "Sorry, I'm having trouble accessing the database.").
- **Ollama Integration:** We use the community `langchain-ollama` provider to call an Ollama LLM (running locally or self-hosted). If the primary API fails, one fallback is to call Ollama's endpoint. Conversely, if Ollama fails (e.g. returns incoherent output), we fallback to OpenAI.

These patterns reflect LangChain best practices of modular agents, tool registration, and robust fallbacks ¹⁵

¹³ .

SECURITY.md

This document outlines security controls, secret handling, and compliance mapping.

- **DevSecOps Practices:** We shift security left by automating scans in CI. For example, we use pre-commit hooks (e.g. `git-secrets` , `truffleHog`) to prevent committing credentials, and regular GitHub secret scanning ¹⁶ ¹⁷ . The OWASP DevSecOps guideline recommends pre-commit and CI scans so that secrets never enter the repo ¹⁶ .
- **Secrets Management:**
 - *GitHub Actions:* All API keys and tokens (Cloudflare API token, OpenAI key, Pinecone key, Notion token) are stored as encrypted GitHub Actions secrets, not in code. GitHub encrypts these and never logs their values ¹⁷ .
 - *Cloudflare Workers:* We use Cloudflare's **Secrets** store (via `wrangler secrets`) for values used in Workers ⁴ . Local development uses `.env` or `.dev.vars` files (see Cloudflare docs) with those secrets, and **never commits them** to Git ⁴ . These files are gitignored.
- **Observability:** We integrate Sentry for error monitoring. The build pipeline uploads source maps (by setting `upload_source_maps = true` in `wrangler.toml`) so stack traces are deminified ¹⁸ . All unhandled exceptions in Workers and frontend code are captured by Sentry, providing alerts and context for investigations.
- **OWASP Top 10:** The code is reviewed against OWASP Top 10 risks. For example, we validate and sanitize all user inputs to prevent injection/prompt-injection attacks ¹⁹ . Authentication and authorization checks prevent broken access control. HTTPS is enforced and API requests are rate-limited to mitigate DoS.
- **SOC 2 & GDPR:** We map our controls to compliance standards:
 - *SOC 2:* We document system descriptions and control matrices, covering change management (CI/CD audits), incident response (Sentry logs), and confidentiality (encrypted secrets). All production data access is logged and monitored to satisfy SOC 2 requirements.
 - *GDPR:* Personal data (if any in Notion or other DBs) is identified and processed lawfully. We provide data deletion flows and use data minimization. Cloudflare's EU data residency options

and our own policies ensure privacy controls. (A full data flow diagram is maintained separately for audit purposes.)

In summary, we follow secure coding and deployment guidelines (e.g. OWASP standards ¹⁹), use encrypted secret stores ⁴ ¹⁷, and implement monitoring for rapid incident response.

API.md

This documents each Next.js API endpoint (backend route handler) with its request/response schema, auth, and examples.

- **Endpoint Reference:** For each route (e.g. `POST /api/chat`, `GET /api/data`), list:
 - URL path and HTTP method.
 - Purpose description.
 - Required headers (e.g. `Authorization: Bearer <JWT>` or API key).
 - Rate limits (e.g. "100 requests/min per user").
 - Request payload schema (JSON fields and types).
 - Response format (JSON structure, status codes).
 - Example request/response body.
- **Example:**
 - `POST /api/chat`: body `{ "messages": [...], "userId": "123" }`. Auth via session cookie or JWT. Returns `{ "reply": "...", "tokenUsage": 56 }`.
 - `GET /api/index`: triggers Pinecone index rebuild. No payload.
- **Auth & Error Codes:** Specify which endpoints require authentication and what happens on invalid access. List HTTP status codes and error JSON payloads. For example, `401 Unauthorized` if no token, or `429 Too Many Requests` if rate-limited.
- **Data Modeling:** For endpoints interacting with Pinecone or Notion, include how we structure documents.
- **Interactive Docs:** Encourage using OpenAPI schemas or Markdown tables. Good API docs include clear endpoint details, parameters, request/response examples, auth methods, and error handling ²⁰.
- **Next.js Routing:** Note that Next.js maps `pages/api/*.js` files to routes ²¹. For example, `pages/api/chat.js` becomes `/api/chat`.
- **Usage Limits:** Document any quotas (e.g. Pinecone requests/min, Ollama query limits) so clients know to handle rate limiting.

Proper API docs follow industry guidelines: precise endpoint details, example payloads, and authentication info ²⁰ ²¹.

TESTING.md

This file describes our testing strategy and coverage goals.

- **Unit Testing:** We write unit tests (using e.g. Jest or PyTest) for individual functions/chains. Adopt TDD where possible. Use mocking to isolate components (e.g. mock OpenAI, Pinecone clients) ²². Aim for high coverage, especially for complex logic. Parameterized tests cover edge cases.
- **Integration Testing:** Tests that span multiple components. For example, a test might spin up a development server and test API routes end-to-end without the UI. We test integration with external APIs (mocking third-party responses) and databases. Key flows (e.g. "vector search returns relevant items") are validated ²³.

- **End-to-End (E2E) Testing:** We run E2E tests simulating real user flows (e.g. using Playwright or Cypress for UI, or full HTTP requests for APIs) to ensure the system works as a whole ²³ ²⁴ . This includes performance tests (under load) and security tests (attempt known attack payloads). Synthetic monitoring also serves as scheduled E2E checks in production.
- **Resilience Tests:** We perform fault-injection tests to exercise circuit-breakers and retries. For instance, we simulate an OpenAI API outage and verify that the fallback (Ollama) is called and the user still gets a valid response. This ensures our fallback strategies are tested.
- **Coverage and Automation:** All tests are automated in CI. We enforce coverage thresholds. Unit tests run on every push; integration/E2E tests run nightly and on release branches. Results, including failed tests and coverage reports, are published to the team dashboard.
- **Test Data:** Use representative datasets (or synthetic data) for testing vector search and prompt responses. Isolate test database from prod. Use `Docker` or in-memory databases for CI.
- **Monitoring in Tests:** We track synthetic monitoring results and pipeline health (e.g. via health-check endpoints) to alert on failures automatically.

These practices align with recommended testing: mock dependencies for unit tests, integration tests across components, and comprehensive E2E flows ²² ²³ . We also include chaos/circuit-breaker tests to validate fault tolerance.

CI_CD.md

This describes our automated pipeline in GitHub Actions and Cloudflare builds.

- **Pipeline Overview:** On each commit to `main`, GitHub Actions triggers the pipeline. It performs linting, unit tests, builds the project, uploads assets, and deploys to Cloudflare Workers/Pages. Cloudflare's own Workers Builds could alternatively be used, but we use GitHub Actions for flexibility ⁸ .
- **Jobs and Steps:** A typical workflow:
 - **Checkout & Setup:** Pull code, setup Node/LangChain env.
 - **Lint & Static Analysis:** Run ESLint/TSLint and any type checks.
 - **Unit Tests:** Run tests and report coverage.
 - **Build:** Compile the Next.js app (including source maps). For Workers, set `upload_source_maps = true` in `wrangler.toml` ¹⁸ .
 - **Sentry Integration:** Use Sentry CLI or GitHub Action to upload generated source maps, enabling readable stack traces in Sentry ¹⁸ .
 - **Deploy:** Use the official Cloudflare Wrangler Action to run `wrangler deploy` with the API token and account ID from secrets ⁷ . (Secrets `CLOUDFLARE_API_TOKEN` and `CLOUDFLARE_ACCOUNT_ID` are configured in GitHub without exposing them in code ⁷ .)
 - **Health Check:** After deployment, a step pings a health-check endpoint (e.g. `/api/health`) to verify the app is responding. Failures here abort the pipeline.
 - **Key Rotation (Maintenance):** We include a scheduled job or script that can rotate API keys/tokens (e.g. daily or weekly). This uses the GitHub Actions and Cloudflare APIs to refresh credentials automatically.

flowchart LR

```

A[Code Pushed] --> B[GitHub Actions CI]
B --> C[Lint & Static Analysis]
C --> D[Run Unit Tests]
D --> E[Build App & Source Maps]
E --> F[Upload Source Maps to Sentry]

```

```

F --> G[Wrangler Deploy to Cloudflare]
G --> H[Post-Deploy Health Check]
H --> I{Success?}
I -- No --> J[Alert/Metrics]
I -- Yes --> K[Done]

```

- **CI/CD Best Practices:** According to Cloudflare, automating deploys ensures consistency and security ⁸. We store tokens as secrets (see above) and use lock-step deployments. Sentry integration is done in CI by configuring the `upload_source_maps` flag and using Sentry's wizard or CLI ¹⁸.
- **Branching:** Deploys happen on `main`. Feature branches can also build to preview environments using Cloudflare Preview.
- **Monitoring:** The pipeline reports to team chat and dashboard. Failed steps (lint, tests, or health check) block merging.
- **Documentation:** The CI workflow file (`.github/workflows/deploy.yml`) includes comments and references. We also version the Terraform/Wrangler config for infra-as-code.

This CI/CD flow follows Cloudflare's and industry guidelines: automated builds/deploys for consistent releases ⁸, secure secret management, and health checks post-deploy.

Sources: This structure is informed by engineering documentation standards and industry best practices ¹ ² ⁶ ¹⁹ ²⁰ ⁸, as well as LangChain and Cloudflare technical guides ¹³ ¹⁸ ⁷. All sections are written in clear, structured Markdown with diagrams (Mermaid and embedded images) to aid understanding.

¹ Documentation Best Practices | styleguide

https://google.github.io/styleguide/docguide/best_practices.html

² ³ ⁵ How to write README files like a pro - DEV Community

<https://dev.to/grski/readme-file-structure-and-sections-pm8>

⁴ Secrets · Cloudflare Workers docs

<https://developers.cloudflare.com/workers/configuration/secrets/>

⁶ Mermaid Chart - Create complex, visual diagrams with text. A smarter way of creating diagrams.

<https://docs.mermaidchart.com/mermaid-oss/syntax/architecture.html>

⁷ GitHub Actions · Cloudflare Workers docs

<https://developers.cloudflare.com/workers/ci-cd/external-cicd/github-actions/>

⁸ CI/CD · Cloudflare Workers docs

<https://developers.cloudflare.com/workers/ci-cd/>

⁹ ¹⁰ LangChain & Multi-Agent AI in 2025: Framework, Tools & Use Cases

<https://blogs.infoservices.com/artificial-intelligence/langchain-multi-agent-ai-framework-2025/>

¹¹ ¹² Integration with MCP Smithery - Feature Requests - Cursor - Community Forum

<https://forum.cursor.com/t/integration-with-mcp-smithery/57005>

¹³ Fallbacks | Langchain

https://js.langchain.com/docs/how_to/fallbacks/

¹⁴ How to handle tool errors | LangChain

https://python.langchain.com/docs/how_to/tools_error/

15 19 22 23 24 **LangChain Development Best Practices | Cursor Rules Guide | cursorrules**
<https://cursorrules.org/article/langchain-cursor-mdc-file>

16 **OWASP DevSecOps Guideline - v-0.2 | OWASP Foundation**
<https://owasp.org/www-project-devsecops-guideline/latest/01a-Secrets-Management>

17 **Best practices for GitHub Action secrets management**
<https://graphite.dev/guides/best-practices-for-github-action-secrets-management>

18 **Source Maps | Sentry for Cloudflare**
<https://docs.sentry.io/platforms/javascript/guides/cloudflare/sourcemaps/>

20 **The Ultimate Guide to API Documentation Best Practices (2025 Edition) | Theneo Blog**
<https://www.theneo.io/blog/api-documentation-best-practices-guide-2025>

21 **Routing: API Routes | Next.js**
<https://nextjs.org/docs/pages/building-your-application/routing/api-routes>